Artificial Intelligence EDAP01

Lecture 4: Machine Learning (I)

Pierre Nugues

Lund University
Pierre.Nugues@cs.lth.se
http://cs.lth.se/pierre_nugues/

January 29, 2021

Why Machine Learning: Early Artificial Intelligence

Early artificial intelligence techniques used introspection to codify knowledge, often in the form of rules.

Expert systems, one of the most notable applications of traditional AI, were entirely based on the competence of experts.

This has two major drawbacks:

- Need of an expertise to understand and explain the rules
- Bias introduced by the expert

Why Machine Learning: What has Changed

Now terabytes of data available.

Makes it impossible to understand such volumes of data, organize them using manually-crafted rules.

Triggered a major move to empirical and statistical techniques.

Applications in natural language processing, medicine, banking, online shopping, image recognition, etc.

> *The success of companies like Google, Facebook, Amazon, and Netflix, not to mention Wall Street firms and industries from manufacturing and retail to healthcare, is increasingly driven by better tools for extracting meaning from very large quantities of data. 'Data Scientist' is now the hottest job title in Silicon Valley.*

– Tim O'Reilly

Some Definitions

1. Machine learning always starts with **datasets**: a collection of objects or observations.

2. Machine-learning algorithms can be classified along two main lines: **supervised** and **unsupervised** classification.

3. Supervised algorithms need a **training set**, where the objects are described in terms of attributes and belong to a known class or have a known output.

4. The performance of the resulting classifier is measured against a **test set**.

5. We can also use N-fold cross validation, where the test set is selected randomly from the training set N times, usually 10.

6. Unsupervised algorithms consider objects, where no class is provided.

7. Unsupervised algorithms learn regularities in datasets.

A Dataset: *Salammbô*

A corpus is a collection – a body – of texts.

French original	English translation

Supervised Learning

Letter counts from *Salammbô*

Chapter	French		English	
	# characters	# A	# characters	# A
Chapter 1	36,961	2,503	35,680	2,217
Chapter 2	43,621	2,992	42,514	2,761
Chapter 3	15,694	1,042	15,162	990
Chapter 4	36,231	2,487	35,298	2,274
Chapter 5	29,945	2,014	29,800	1,865
Chapter 6	40,588	2,805	40,255	2,606
Chapter 7	75,255	5,062	74,532	4,805
Chapter 8	37,709	2,643	37,464	2,396
Chapter 9	30,899	2,126	31,030	1,993
Chapter 10	25,486	1,784	24,843	1,627
Chapter 11	37,497	2,641	36,172	2,375
Chapter 12	40,398	2,766	39,552	2,560
Chapter 13	74,105	5,047	72,545	4,597
Chapter 14	76,725	5,312	75,352	4,871
Chapter 15	18,317	1,215	18,031	1,119

Dataset: https://github.com/pnugues/ilppp/tree/master/
programs/ch04/salammbo

Supervised Learning: Regression

Letter count from *Salammbô* in French
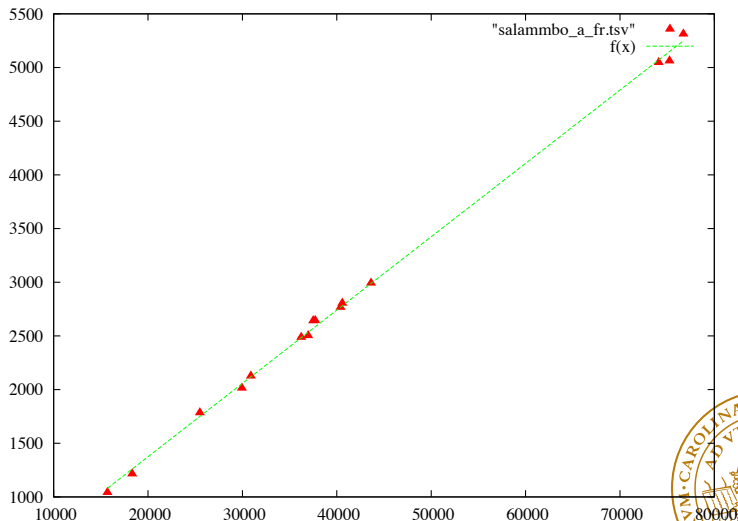
Models

We will assume that datasets are governed by functions or models.
For instance given the set:

$$\{(\mathbf{x}_i, y_i) | 0 < i \leqslant N\},$$

there exists a function such that:

$$f(\mathbf{x}_i) = y_i.$$

Supervised machine learning algorithms will produce hypothesized functions or models fitting the data.
The predicted value is denoted:

$$f(\mathbf{x}_i) = \hat{y}_i.$$

to make a difference with the observed value: y_i.

Selecting a Model

Often, multiple models can fit a dataset:
Three polynomials of degree: 1, 8, and 9

A general rule in machine learning is to prefer the simplest hypotheses, here the lower polynomial degrees.

Source code:
https://github.com/pnugues/ilppp/tree/master/programs/ch04

Supervised Learning: Classification

A dataset for binary classification.

	# char.	# A	class (y)	# char.	# A	class (y)
Chapter 1	36,961	2,503	1	35,680	2,217	0
Chapter 2	43,621	2,992	1	42,514	2,761	0
Chapter 3	15,694	1,042	1	15,162	990	0
Chapter 4	36,231	2,487	1	35,298	2,274	0
Chapter 5	29,945	2,014	1	29,800	1,865	0
Chapter 6	40,588	2,805	1	40,255	2,606	0
Chapter 7	75,255	5,062	1	74,532	4,805	0
Chapter 8	37,709	2,643	1	37,464	2,396	0
Chapter 9	30,899	2,126	1	31,030	1,993	0
Chapter 10	25,486	1,784	1	24,843	1,627	0
Chapter 11	37,497	2,641	1	36,172	2,375	0
Chapter 12	40,398	2,766	1	39,552	2,560	0
Chapter 13	74,105	5,047	1	72,545	4,597	0
Chapter 14	76,725	5,312	1	75,352	4,871	0
Chapter 15	18,317	1,215	1	18,031	1,119	0

Supervised Learning: Regression and Classification

Given the dataset, $\{(\mathbf{x}_i, y_i) | 0 < i \leqslant N\}$ and a model f:

- Classification: $f(\mathbf{x}) = y$ is discrete,
- Regression: $f(\mathbf{x}) = y$ is continuous.

Types of Iris

Iris virginica   Iris setosa   Iris versicolor

Courtesy Wikipedia

Supervised Learning: Fisher's Iris dataset (1936)

180 MULTIPLE MEASUREMENTS IN TAXONOMIC PROBLEMS

Table I

Iris setosa				Iris versicolor				Iris virginica			
Sepal length	Sepal width	Petal length	Petal width	Sepal length	Sepal width	Petal length	Petal width	Sepal length	Sepal width	Petal length	Petal width
5·1	3·5	1·4	0·2	7·0	3·2	4·7	1·4	6·3	3·3	6·0	2·5
4·9	3·0	1·4	0·2	6·4	3·2	4·5	1·5	5·8	2·7	5·1	1·9
4·7	3·2	1·3	0·2	6·9	3·1	4·9	1·5	7·1	3·0	5·9	2·1
4·6	3·1	1·5	0·2	5·5	2·3	4·0	1·3	6·3	2·9	5·6	1·8
5·0	3·6	1·4	0·2	6·5	2·8	4·6	1·5	6·5	3·0	5·8	2·2
5·4	3·9	1·7	0·4	5·7	2·8	4·5	1·3	7·6	3·0	6·6	2·1
4·6	3·4	1·4	0·3	6·3	3·3	4·7	1·6	4·9	2·5	4·5	1·7
5·0	3·4	1·5	0·2	4·9	2·4	3·3	1·0	7·3	2·9	6·3	1·8
4·4	2·9	1·4	0·2	6·6	2·9	4·6	1·3	6·7	2·5	5·8	1·8
4·9	3·1	1·5	0·1	5·2	2·7	3·9	1·4	7·2	3·6	6·1	2·5
5·4	3·7	1·5	0·2	5·0	2·0	3·5	1·0	6·5	3·2	5·1	2·0
4·8	3·4	1·6	0·2	5·9	3·0	4·2	1·5	6·4	2·7	5·3	1·9
4·8	3·0	1·4	0·1	6·0	2·2	4·0	1·0	6·8	3·0	5·5	2·1
4·3	3·0	1·1	0·1	6·1	2·9	4·7	1·4	5·7	2·5	5·0	2·0
5·8	4·0	1·2	0·2	5·6	2·9	3·6	1·3	5·8	2·8	5·1	2·4
5·7	4·4	1·5	0·4	6·7	3·1	4·4	1·4	6·4	3·2	5·3	2·3
5·4	3·9	1·3	0·4	5·6	3·0	4·5	1·5	6·5	3·0	5·5	1·8
5·1	3·5	1·4	0·3	5·8	2·7	4·1	1·0	7·7	3·8	6·7	2·2
5·7	3·8	1·7	0·3	6·2	2·2	4·5	1·5	7·7	2·6	6·9	2·3
5·1	3·8	1·5	0·3	5·6	2·5	3·9	1·1	6·0	2·2	5·0	1·5
5·4	3·4	1·7	0·2	5·9	3·2	4·8	1·8	6·9	3·2	5·7	2·3
5·1	3·7	1·5	0·4	6·1	2·8	4·0	1·3	5·6	2·8	4·9	2·0

Berkson's Dataset (1944)

Drug concentration	Number exposed	Survive Class 0	Die Class 1	Mortality rate	Expected mortality
40	462	352	110	.2359	.2206
60	500	301	199	.3980	.4339
80	467	169	298	.6380	.6085
100	515	145	370	.7184	.7291
120	561	102	459	.8182	.8081
140	469	69	400	.8529	.8601
160	550	55	495	.9000	.8952
180	542	43	499	.9207	.9195
200	479	29	450	.9395	.9366
250	497	21	476	.9577	.9624
300	453	11	442	.9757	.9756

Table: A dataset. Adapted and simplified from the original article that described how to apply logistic regression to classification by Joseph Berkson, Application of the Logistic Function to Bio-Assay. *Journal of the American Statistical Association* (1944).

Regression: Another Model, the Logistic Curve (Verhulst)

Unsupervised Learning: Clustering

No class is given:

Objects, Attributes, and Classes. After Quinlan (1986)

Object	Attributes				Class
	Outlook	Temperature	Humidity	Windy	
1	Sunny	Hot	High	False	N
2	Sunny	Hot	High	True	N
3	Overcast	Hot	High	False	P
4	Rain	Mild	High	False	P
5	Rain	Cool	Normal	False	P
6	Rain	Cool	Normal	True	N
7	Overcast	Cool	Normal	True	P
8	Sunny	Mild	High	False	N
9	Sunny	Cool	Normal	False	P
10	Rain	Mild	Normal	False	P
11	Sunny	Mild	Normal	True	P
12	Overcast	Mild	High	True	P
13	Overcast	Hot	Normal	False	P
14	Rain	Mild	High	True	N

Classifying Objects with Decision Trees. After Quinlan (1986)

Matrix Notation

- A feature vector (predictors or attributes): $\mathbf{x}$, and feature matrix: $\mathbf{X}$;
- The class: y and the class vector: $\mathbf{y}$;
- The predicted class (response): $\hat{y}$, and predicted class vector: $\hat{\mathbf{y}}$

$$\mathbf{X} = \begin{bmatrix} Sunny & Hot & High & False \\ Sunny & Hot & High & True \\ Overcast & Hot & High & False \\ Rain & Mild & High & False \\ Rain & Cool & Normal & False \\ Rain & Cool & Normal & True \\ Overcast & Cool & Normal & True \\ Sunny & Mild & High & False \\ Sunny & Cool & Normal & False \\ Rain & Mild & Normal & False \\ Sunny & Mild & Normal & True \\ Overcast & Mild & High & True \\ Overcast & Hot & Normal & False \\ Rain & Mild & High & True \end{bmatrix} ; \mathbf{y} = \begin{bmatrix} N \\ N \\ P \\ P \\ P \\ N \\ P \\ N \\ P \\ P \\ P \\ P \\ P \\ N \end{bmatrix}$$

Decision Trees and Classification

- Each object is defined by an attribute vector (or feature vector or predictors) $\{A_1, A_2, ..., A_v\}$
- Each object belongs to one class $\{C_1, C_2, ..., C_n\}$
- The attributes of the examples are: $\{Outlook, Temperature, Humidity, Windy\}$ and the classes are: $\{N, P\}$.
- The nodes of the tree are the attributes.
- Each attribute has a set of possible values. The values of *Outlook* are $\{sunny, rain, overcast\}$
- The branches correspond to the values of each attribute
- The optimal tree corresponds to a minimal number of tests

Entropy, Decision Trees, and Classification

- Decision trees are useful devices to classify objects into a set of classes.
- Shannon's entropy is a quantification of distribution diversity in a set of symbols
- It can help us learn automatically decision trees from a set of data.
- The algorithm is one of the simplest machine-learning techniques to obtain a classifier.

Inducing (Learning) Decision Trees Automatically: ID3

It is possible to design many trees that classify the objects successfully
An efficient decision tree uses a minimal number of tests.
In the decision tree:

- Each example is defined by a finite number of attributes.
- Each node in the decision tree corresponds to an attribute that has as many branches as the attribute has possible values.

At the root of the tree, the condition must be the most discriminating, that is, have branches gathering most positive examples while others gather negative examples.

ID3 (Quinlan 1986)

Each attribute scatters the set into as many subsets as there are values for this attribute.

In the dataset, we have nine positive examples and five negative ones,

At each decision point, the "best" attribute has the maximal separation power, the maximal information gain

Better Root?

```
          ┌─────────────────┐
          │   Temperature   │
          │   P: 9, N: 5    │
          └─────────────────┘
         hot      mild       cool
   ┌──────────┐ ┌──────────┐ ┌──────────┐
   │ P: 2, N: 2│ │ P: 4, N: 2│ │ P: 3, N: 1│
   └──────────┘ └──────────┘ └──────────┘
```

How can we measure this?
Answer: Entropy or Gini impurity

Entropy (I)

Information theory models a text as a sequence of symbols.
Let $x_1, x_2, ..., x_N$ be a discrete set of N symbols representing the characters.
The **information content** of a symbol is defined as

$$I(x_i) = -\log_2 P(x_i) = \log_2 \frac{1}{P(x_i)},$$

where

$$P(x_i) = \frac{Count(x_i)}{\sum_{j=1}^{n} Count(x_j)}.$$

Information content of:

- P: $-\log_2 \frac{9}{14} = 0.6374$
- N: $-\log_2 \frac{5}{14} = 1.4854$

Entropy (II)

Entropy, the average information content, is defined as:

$$H(X) = - \sum_{x \in X} P(x) \log_2 P(x),$$

By convention: $0 \log_2 0 = 0$.

Entropy of the dataset:

$$I(P, N) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940.$$

Entropy Examples

Outlook
P: 9, N: 5

sunny   overcast   rain

P: 2, N: 3 P: 4 P: 3, N: 2

In the three child nodes:

$$
\begin{aligned}
\textit{sunny} : \quad & I(P_1, N_1) &=& \quad -\frac{2}{5}\log_2\frac{2}{5} - \frac{3}{5}\log_2\frac{3}{5} = 0.971. \\
\textit{overcast} : \quad & I(P_2, N_2) &=& \quad 0.0 \\
\textit{rain} : \quad & I(P_3, N_3) &=& \quad -\frac{3}{5}\log_2\frac{3}{5} - \frac{2}{5}\log_2\frac{2}{5} = 0.971
\end{aligned}
$$

Understanding the Entropy

For a two-class set, we set:

$$x = \frac{p}{p+n} \text{ and } \frac{n}{p+n} = 1 - x.$$

$$I(x) = -x \log_2 x - (1-x) \log_2 (1-x) \text{ with } x \in [0, 1].$$

The entropy reaches a maximum when there are as many positive as negative examples in the dataset. It is minimal when the set consists of either positive or negative examples.

Gini Impurity

Gini impurity is an alternative to entropy:

$$G(X) = -\sum_{x \in X} P(x)(1 - P(x)),$$

For a two-class set, we set:

$$
\begin{aligned}
2G(x) &= 2(x(1-x) + (1-x)x) \text{ with } x \in [0, 1]. \\
&= 4x(1-x)
\end{aligned}
$$

Entropy of a Text

The entropy of the text is

$$
\begin{aligned}
H(X) &= -\sum_{x \in X} P(x) \log_2 P(x). \\
&= -P(A) \log_2 P(A) - P(B) \log_2 P(B) - \ldots \\
&\quad -P(Z) \log_2 P(Z) - P(\grave{A}) \log_2 P(\grave{A}) - \ldots \\
&\quad -P(\ddot{Y}) \log_2 P(\ddot{Y}) - P(blanks) \log_2 P(blanks).
\end{aligned}
$$

Entropy of Gustave Flaubert's *Salammbô* in French is $H(X) = 4.39$.

Cross-Entropy

The cross entropy of M on P is defined as:

$$H(P, M) = - \sum_{x \in X} P(x) \log_2 M(x),$$

where typically:

- The model distribution M is learned from a training set;
- The P distribution is obtained from a test set.

We have the inequality $H(P) \leq H(P, M)$.

The difference is called the **Kullback-Leibler divergence**.

Example of Kullback-Leibler divergence.

- *Salammbô*, chapters 1-14, text in French, the model
- *Salammbô*, chapters 15, text in French, a test set
- *Notre Dame de Paris*: text in French, a test set
- *Nineteen Eighty-Four*: text in English, a test set

		Entropy	Cross entr.	Diff.
M	*Salammbô*, chapters 1-14, training set	4.39481	4.39481	0.0
P	*Salammbô*, chapter 15, test set	4.34937	4.36074	0.01137
P	*Notre Dame de Paris*, test set	4.43696	4.45507	0.01811
P	*Nineteen Eighty-Four*, test set	4.35922	4.82012	0.46090

ID3 (Quinlan 1986)

ID3 uses the entropy to select the best attribute to be the root of the tree and recursively the next attributes of the resulting nodes.
The entropy of a two-class set p and n is:

$$I(p, n) = -\frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}.$$

The weighted average of all the nodes below an attribute A is:

$$E(A) = \sum_{i=1}^{v} \frac{p_i + n_i}{p+n} I(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}).$$

The information gain is defined as $I(p, n) - E(A)$ (or $I_{before} - I_{after}$).
This measures the separating power of an attribute: the more the gain, the better the attribute.

Example

The entropy of the dataset is:

$$I(p, n) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940.$$

Outlook has three values: *sunny*, *overcast*, and *rain*. The entropies of the corresponding subsets are:

$$
\begin{array}{llcl}
sunny: & I(p_1, n_1) & = & -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} = 0.971. \\
overcast: & I(p_2, n_2) & = & 0.0 \\
rain: & I(p_3, n_3) & = & -\frac{3}{5} \log_2 \frac{3}{5} - \frac{2}{5} \log_2 \frac{2}{5} = 0.971.
\end{array}
$$

Gain(*Outlook*) is then $0.940 - \frac{5}{14} \times 0.971 - \frac{5}{14} \times 0.971 = 0.246$, the highest possible among the attributes.
We have *Gain*(*Temperature*) = 0.029, *Gain*(*Humidity*) = 0.151, and *Gain*(*Windy*) = 0.048.

Inducing Decision Trees Automatically: ID3

The algorithm to build the decision tree is simple:

1. The information gain is computed on the dataset for all attributes in the set of attributes, *SA*.

2. The attribute with the highest gain is selected to be the root of the tree:
$$A = \arg\max_{a \in SA} I(n, p) - E(a).$$

3. The dataset is split into v subsets $\{N_1, ..., N_v\}$, where the value of A for the objects in N_i is A_i

4. For each subset, a corresponding node is created below the root.

5. This process is repeated recursively for each node of the tree with the subset it contains until all the objects of the node are either positive or negative.

For a training set of N instances each having M attributes, ID3's complexity to generate a decision tree is $O(NM)$.

# Decision Tree Learning Algorithm (From the Textbook)

1: **function** DTL(*examples*, *attributes*, *parent_examples*) **returns** a tree
2:     **if** *examples* is empty **then**
3:         **return** PluralityValue(*parent_examples*)
4:     **else if** all *examples* have the same classification **then**
5:         **return** the classification
6:     **else if** *attributes* is empty **then**
7:         **return** PluralityValue(*examples*)
8:     **else**
9:         $A \leftarrow \arg\max_{a \in attributes} \text{InformationGain}(a, examples)$
10:         *tree* ← a new decision tree with root test A
11:         **for all** $v_k \in A$ **do**
12:             *exs* ← $\{e : e \in examples \text{ and } e.A = v_k\}$
13:             *subtree* ← DTL(*exs*, *attributes* − *A*, *examples*)
14:             add a branch to *tree* with label $\{A = v_k\}$ and subtree
    subtree **return** *tree*

Learning Curve

The classical evaluation technique uses a training set and a test set. Generally, the larger the training set, the better the performance. This can be visualized with a learning curve. From the textbook, Stuart Russell and Peter Norvig, *Artificial Intelligence*, 3rd ed., 2010, page 703.

Overfitting

- When two classifiers have equal performances on a specific test set, the simplest one is supposed to be more general
- A small decision tree is always preferable to a larger one.
- Complex classifiers may show an overfit to the training data and have poor performance when the dataset changes.
- In the case of decision trees, overfits show with deep trees and when numerous leaves have few examples.

Pruning Trees

Decision tree pruning consists in removing and merging leaves with few objects, for instance one or two.

To prune the tree:

1. Generate the complete decision tree
2. Consider the nodes that have only leaves as children and merge the leaves if the information gain is below a certain threshold.
3. To determine the threshold, we use a significance test.

This is the χ^2 pruning.

Contingency Tables

Example: A set of 12 positive and 8 negative examples.
Let us compare two attributes:

- Attribute 1 with three values: $\{v_1^1, v_1^2, v_1^3\}$
- Attribute 2 with three values: $\{v_2^1, v_2^2, v_2^3\}$

Contingency tables are devices to visualize frequency distributions.

Attribute 1		v_1^1	v_1^2	v_1^3	
	Positive	6	3	3	12 (60%)
	Negative	4	2	2	8 (40%)
	Total	10	5	5	20 (100%)

Attribute 2		v_2^1	v_2^2	v_2^3	
	Positive	1	1	10	12 (60%)
	Negative	8	0	0	8 (40%)
	Total	9	1	10	20 (100%)

Significance Test

An attribute with no significance would have the same proportions of positive and negative examples before and after the test.

Proportions before, for instance 12 and 8:

$$\frac{p}{p+n} \quad \text{and} \quad \frac{n}{p+n}$$

Proportions after, for v_k, a value of the attribute:

$$\frac{p_k}{p_k+n_k} \quad \text{and} \quad \frac{n_k}{p_k+n_k}$$

The expected values of a nonsignificant attribute (null hypothesis) are:

$$\hat{p}_k = \frac{p}{p+n} \times (p_k+n_k) \qquad \hat{n}_k = \frac{n}{p+n} \times (p_k+n_k)$$

Example

Contingency tables with the expected frequencies, in red.

Attribute 1		v_1^1	v_1^1	v_1^2	v_1^2	v_1^3	v_1^3	
	Positive	6	6	3	3	3	3	12 (60%)
	Negative	4	4	2	2	2	2	8 (40%)
	Total	10		5		5		20 (100%)

Attribute 2		v_1^1	v_1^1	v_1^2	v_1^2	v_1^3	v_1^3	
	Positive	1	5.4	1	0.6	10	6	12 (60%)
	Negative	8	3.6	0	0.4	0	4	8 (40%)
	Total	9		1		10		20 (100%)

The total deviation for an attribute with values $v_1..v_d$ is:

$$\Delta = \sum_{k=1}^{d} \frac{(p_k - \hat{p}_k)^2}{\hat{p}_k} + \frac{(n_k - \hat{n}_k)^2}{\hat{n}_k}$$

Pruning Trees (II)

- The null hypothesis corresponds to an irrelevant attribute: $\Delta = 0$.
- Δ values are distributed according to the χ^2 distribution
- If v is the number of values of an attribute, the degree of freedom is $v - 1$
- The corresponding χ^2 distributions can be found in statistical tables.

  - For three degrees of freedom:
    10%: 6.25, 5%: 7.81, 1%: 11.34, 1‰: 16.27
  - The significance level is usually 5%

- An alternate pruning method is to modify the termination condition of the algorithm:

  - We stop when the information gain is below a certain threshold

Dealing with Real Data

- Unknown attribute values (imputation): Use the most frequent value of the attribute or all the values;

- Attributes with many values (large attribute domains): Difficult to handle for decision trees. The book suggests to use the gain ratio:

$$\frac{\text{Information gain}}{\text{Information value}}, \text{where Information value} = -\sum_{i=1}^{v} \frac{p_i + n_i}{p + n} \log_2 \frac{p_i + n_i}{p + n}$$

- Numerical attributes: Find the binary split point that maximizes the information gain

- Numerical output: Regression tree

Evaluation

- The standard evaluation procedure is to train the classifier on a training set and evaluate the performance on a test set.
- When we have only one set, we divide it in two subsets: the training set and the test set (or holdout data).
- The split can be 90–10 or 80–20
- This often optimizes the classifier for a specific test set and creates an overfit

A possible solution is to have a validation (or development) set and carry out intermediate evaluations on it.

Cross Validation

- A *N*-fold cross validation mitigates the overfit
- The set is partitioned into *N* subsets, $N = 5$ for example, one of them being the test set (red) and the rest the training set (gray).
- The process is repeated *N* times with a different test set: *N* folds

At the extreme, **leave-one-out cross-validation**.

Model Selection

- Validation can apply to one classification method
- We can use it to select a classification method and its parametrization.
- Needs three sets: training set, development set, and test set.

Loss

- Depending on their type, errors can have different impacts.
- Imagine a smoke detector that predicts fire.
- Equivalent to a classifier that uses input data from sensors and classifies them as *fire* and *nonfire*.
- The loss function is defined as $L(y, \hat{y})$ with $\hat{y} = h(\mathbf{x})$, where $\mathbf{x}$ is the vector of attributes, h the classifier, and y, the correct value.
- The cost of $L(fire, nonfire)$ is much higher than $L(nonfire, fire)$.
- $L(fire, fire) = L(nonfire, nonfire) = 0$

Common Loss Measures

The loss function is defined as $L(y, \hat{y})$ with $\hat{y} = h(\mathbf{x})$, where $\mathbf{x}$ is the vector of attributes, h the classifier, and y, the correct value.

Absolute value loss	$L_1(y, \hat{y})$	$=$	$\lvert y - \hat{y} \rvert$
Squared error loss	$L_2(y, \hat{y})$	$=$	$(y - \hat{y})^2$
0/1 loss	$L_{0/1}(y, \hat{y})$	$=$	0 if $y = \hat{y}$ else 1
Binary crossentropy			
Categorical crossentropy			

Empirical Loss

We compute the empirical loss of a classifier h on a set of examples E using the formula:

$$Loss(L, E, h) = \frac{1}{N} \sum_E L(y, h(x)).$$

For continuous functions:

$$Loss(L, E, h) = \frac{1}{N} \sum_E (y - h(x))^2.$$

Evaluation

There are different kinds of measures to evaluate the performance of machine learning techniques, for instance:

- Precision and recall in information retrieval and natural language processing;
- The *receiver operating characteristic* (ROC) in medicine.

	Positive examples: P	Negative examples: N
Classified as P	True positives: A	False positives: B
Classified as N	False negatives: C	True negatives: D

More on the receiver operating characteristic here: `http://en.wikipedia.org/wiki/Receiver_operating_characteristic`

Recall, Precision, and the F-Measure

The **accuracy** is $\frac{|A \cup D|}{|P \cup N|}$.

Recall measures how much relevant examples the system has classified correctly, for P:

$$\text{Recall} = \frac{|A|}{|A \cup C|}.$$

Precision is the accuracy of what has been returned, for P:

$$\text{Precision} = \frac{|A|}{|A \cup B|}.$$

Recall and precision are combined into the **F-measure**, which is defined as the harmonic mean of both numbers:

$$F = \frac{2 \cdot \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}.$$

Measuring Quality: The Confusion Matrix

A task in natural language processing: Identify the parts of speech (POS) of words.

Example: *The can rusted*

- The human: *The*/art (DT) *can*/noun (NN) *rusted*/verb (VBD)
- The POS tagger: *The*/art (DT) *can*/modal (MD) *rusted*/verb (VBD)

↓Correct	Tagger →									
	DT	IN	JJ	NN	RB	RP	VB	VBD	VBG	VBN
DT	99.4	0.3	–	–	0.3	–	–	–	–	–
IN	0.4	97.5	–	–	1.5	0.5	–	–	–	–
JJ	–	0.1	93.9	1.8	0.9	–	0.1	0.1	0.4	1.5
NN	–	–	2.2	95.5	–	–	0.2	–	0.4	–
RB	0.2	2.4	2.2	0.6	93.2	1.2	–	–	–	–
RP	–	24.7	–	1.1	12.6	61.5	–	–	–	–
VB	–	–	0.3	1.4	–	–	96.0	–	–	0.2
VBD	–	–	0.3	–	–	–	–	94.6	–	4.8
VBG	–	–	2.5	4.4	–	–	–	–	93.0	–
VBN	–	–	4.6	–	–	–	–	4.3	–	90.6

After Franz (1996, p. 124)

The Weka Toolkit

Weka: A powerful collection of machine-learning algorithms
http://www.cs.waikato.ac.nz/ml/weka/.

The Weka Toolkit

Running ID3

ARFF: The Weka Data Format

Storing Quinlan's dataset in Weka's attribute-relation file format (ARFF)
http://weka.wikispaces.com/ARFF:

```
@relation weather.symbolic

@attribute outlook {sunny, overcast, rainy}
@attribute temperature {hot, mild, cool}
@attribute humidity {high, normal}
@attribute windy {TRUE, FALSE}
@attribute play {yes, no}

@data
sunny,hot,high,FALSE,no
sunny,hot,high,TRUE,no
overcast,hot,high,FALSE,yes
rainy,mild,high,FALSE,yes
rainy,cool,normal,FALSE,yes
rainy,cool,normal,TRUE,no
overcast,cool,normal,TRUE,yes
```

Other Toolkits

- **Scikit-learn:** https://scikit-learn.org/
- **Keras:** https://keras.io/
- **PyTorch:** https://pytorch.org/